



# WikiNER: Brute-force Named Entity Recognition leveraging *Wikipedia* dataset

COM S 579X Natural Language Processing, Fall 2022  
December 8, 2022

---

*Mario Mastrandrea, Yuting Yang*

# What is Named Entity Recognition?

---

## Text

Iowa State University of Science and  
Technology is a public land-grant  
research university in Ames, Iowa.

## Tags

- [Organization] Iowa State University of Science and Technology
- [Location] Ames
- [Location] Iowa

# What is Named Entity Recognition?

## Text

[Iowa State University of Science and

B I I I I I

Technology]ORG is a public land-grant

I

research university in [Ames]LOC,

B

[Iowa]LOC.

B

## Tagging

## Tags

- [Organization] Iowa State University of Science and Technology
- [Location] Ames
- [Location] Iowa

# What is Named Entity Recognition?

## Text

[Iowa State University of Science and

B I I I I I

Technology]ORG is a public land-grant

I

research university in [Ames]LOC,

B

[Iowa]LOC.

B

## Tagging

## Tags

- [Organization] Entity 1 on *Wikipedia*
- [Location] Entity 2 on *Wikipedia*
- [Location] Entity 2 on *Wikipedia*
- .....


**WikiNER** uses *Wikipedia* entities as Named Entities tags!

## Why WikiNER?

---

The entities in *Wikipedia* are a good source of tags to recognize WHO, WHERE, and WHAT in a sentence

- **Massive** in the number of entities of each type (i.e., organization, locations, persons)
- Being **evolving** and expanding
- Applicable to **different domains**



Use the **top ranked *Wikipedia* Named Entities** to label a sentence with a **brute-force approach**

## How did we build **WikiNER**?

---

1. Retrieve **Wikipedia** top ranked **NEs** from *Wikidata* API
2. Build a **brute-force model** to perform NER tagging
3. **Evaluate** the model using a tagged *corpus* as a benchmark

## Retrieve *Wikipedia* top ranked NEs from *Wikidata* API

1. download **.csv file** from [Wikidata QRank project](#) with **ordered Wikipedia** entities



Entity	QRank
Q178995	219893853
Q635	113674399
Q866	93345399
...	...

2. call **Wikidata API** for each entity to retrieve its **English label**

<https://www.wikidata.org/w/api.php?action=wbgetentities&ids=Q178995&languages=en&format=json&props=labels>



3. create an output **.csv file** with the **top k Named Entities labels**

Entity	QRank	Label
Q178995	219893853	HTTP cookie
Q635	113674399	Cleopatra
Q866	93345399	YouTube
...	...	...

## Retrieve *Wikipedia* top ranked NEs - with aliases

```
def top_N_NEs(  
    input_csv_ranking_file_path: str,  
    output_csv_file_path: str,  
    top_N=1000,  
    aliases=False  
):
```



*The function we implemented to extract the labels*

- It is also possible to retrieve all the **aliases** of a Named Entity with the same API, **increasing the performances!**

<https://www.wikidata.org/w/api.php?action=wbgetentities&ids=Q178995&languages=en&format=json&props=labels|aliases>

Entity	QRank	Label
Q178995	219893853	HTTP cookie
Q178995	219893853	cookie
Q178995	219893853	web cookie
Q178995	219893853	browser cookie
Q178995	219893853	internet cookie
Q178995	219893853	cookies
Q635	113674399	Cleopatra
Q635	113674399	Cleopatra VII Philopator
Q635	113674399	Cleopatra VII
Q866	93345399	YouTube
Q866	93345399	YT
...	...	...



## Build a **brute-force model** to perform **NER tagging**

---

```
def brutal_force_NER(  
    sentence_tokens: list[str],  
    NE_list: list[str],  
    tokenizer,  
    scheme="BIO"  
):
```

*The function we implemented to tag a sentence using the Wikipedia NEs with brute-force approach*

1. Find all the **not-overlapping *Matches*\*** in the sentence using all the NEs (brute-force)
2. Represent this matches according to the specified **tagging scheme** (i.e., *BIO*, *BIOES*)

\*A *Match* is an **occurrence** of a Named Entity in the sentence, represented as a tuple of **indexes**  $(a,b)$ :  $a$  is the **start** index,  $b$  is the (exclusive) **ending** index

## Example of brute-force tagging

sentence\_tokens = ['I', 'went', 'to', 'school', 'at', 'Iowa', 'State', 'University', 'in', 'Fall', '2022']

NE\_list = ['Microsoft', 'Iowa State', 'Fall 2022', 'Deep Learning', 'Iowa State University', '2022', 'school']

(3,4), (5,8), (9,11)



I went to school at Iowa State University in Fall 2022  
O O O B O B I I O B I  
O O O U O B I L O B I

*Matches*

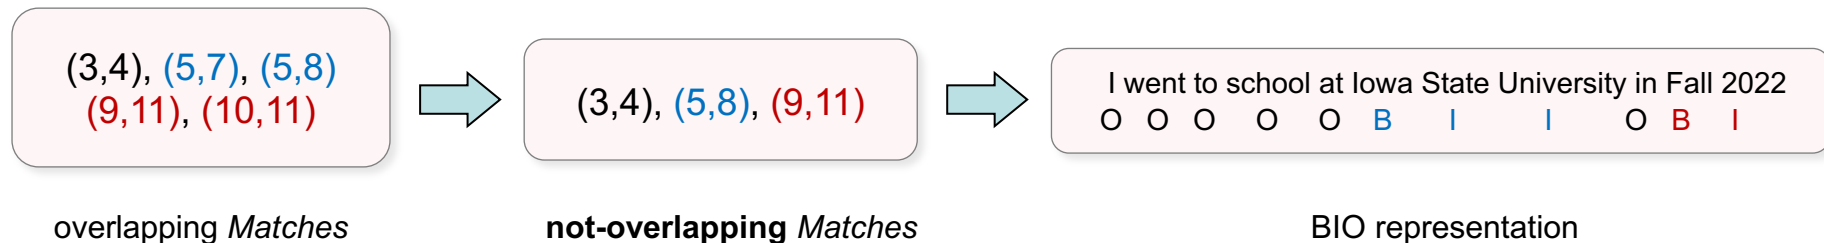
BIO and BIOES representations

## Challenge #1: Overlapping NEs

- We solved the problem of **overlapping** Named Entities prioritizing always the **longest Named Entity**

sentence\_tokens = ['I', 'went', 'to', 'school', 'at', 'Iowa', 'State', 'University', 'in', 'Fall', '2022']

NE\_list = ['Microsoft', 'Iowa State', 'Fall 2022', 'Deep Learning', 'Iowa State University', '2022', 'school']



## Supported tagging schemes

---

- Our **WikiNER** model supports 2 different tagging schemes:

**BIO:**

**B** – Beginning

**I** – Inside

**O** – Outside

**BILOU:**

**B** – Beginning

**I** – Inside

**L** – Last

**O** – Outside

**U** – Unit

## Challenge #2: Entity type (LOC, PER, ORG, MISC, etc.)

---

- NER problem also expects to find the **type** of each Named Entity (LOCation, PERson, ORGanization, etc.)
- But this is **impossible** using just a the *Wikidata* API, since it does **not** provide the type of an entity
- **Solution:** using another model (like **BERT**) to also find the **type** of NEs

## Evaluate the model using a tagged *corpus* as benchmark

- Use **CoNLL2003 dataset** as benchmark
  - It contains a *corpus* of sentences whose Named Entities have been tagged in their respective **types** using the **BIO scheme**
- Use the **WikiNER** model to tag its sentences
- Evaluate the **performances** comparing the **predicted tags** with the **real tags** in the *corpus* (ignoring the NE type)

token	POS tag	chunk tag	NER tag
EU	NNP	B-NP	B-ORG
rejects	VBZ	B-VP	O
German	JJ	B-NP	B-MISC
call	NN	I-NP	O
to	TO	B-VP	O
boycott	VB	I-VP	O
British	JJ	B-NP	B-MISC
lamb	NN	I-NP	O
.	.	O	O
Peter	NNP	B-NP	B-PER
...	...	...	...

## Evaluation metrics

- **Precision**

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

**Among all the predicted NEs, how many of them were **true NEs****

- **Recall**

$$\frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

**Among all the real NEs, how many of them we were able to **predict****

- **F1-score**

$$2 * \frac{\text{Precision} * \text{Recall}}{\text{Precision} + \text{Recall}}$$

**Harmonic mean between *precision* and *recall***

## WikiNER performances on CoNLL2003 dataset

---

- We used **seqeval** library to compute *precision*, *recall* and *F1-score*
- **Compare** the results using **different numbers** of Wikipedia top NEs

#NEs	Precision	Recall	F1-score
100	0.94	0.05	0.09
1,000	0.89	0.14	0.24
10,000	0.26	0.22	0.24

*with **no** aliases*

#NEs	Precision	Recall	F1-score
100	0.82	0.06	0.11
1,000	0.49	0.18	<b>0.26</b>
10,000	0.15	<b>0.29</b>	0.20

*with **aliases***



## Timing problems

---

- *WikiNER* is a very **time-consuming** approach!
  - **extracting** *Wikipedia* NEs
  - **executing** the brute-force model

#NEs	time
100	~20s
1,000	~3min
10,000	~30min

estimated times to **extract** NEs

#NEs	no aliases	with aliases
100	~5min	~20min
1,000	~30min	~2h
10,000	~3h	~12h

estimated times to execute **WikiNER** on **CoNLL2003**

## Use a **BERT Model** to Perform the **NER Tagging**

---

- Due to the **disadvantages** of the brute-force model (i.e., lack of entity type and time-consuming), we used a **BERT** model to perform the task on *CoNLL2003*
- This task takes **four steps**:
  1. Create **tokenized** dataset based on *CoNLL2003* using a BERT tokenizer
  2. Create input objects for **Trainer**, like an **evaluation metrics** function, a **data collator** and the tokenizer
  3. Define **BERT model** and the **training arguments**
  4. **Train** and **evaluate** the BERT model on *CoNLL2003* dataset

## Step 1: Create Tokenized Dataset based on CoNLL2003

- Issue 1: The **BERT model** needs a proper tokenized input, in the form of **integer tokens**, while CoNLL2003 has string tokens
- Solution: **use a BERT tokenizer** to properly **preprocess** input dataset CoNLL2003

```
Tokens:      EU    rejects German call to boycott British lamb .  
NER tag no.: 3     0      7      0      0  0      7      0      0
```



```
str tokens:  [CLS] EU    rejects German call to  boycott British la  ##mb .  [SEP]  
int tokens:  101  7270 22961 1528 1840 1106 21423 1418 2495 12913 119 102
```

example of **tokenization**: 1<sup>st</sup> **sentence** of CoNLL2003 preprocessed using a **BERT tokenizer**

## Step 1: Create Tokenized Dataset – align NER tag labels

- Issue 2: The **Named Entities tags** in *CoNLL2003* need to be formatted in a way that fits the **BERT tokenization**
  - Some words are split in **more tokens**
  - Special tokens** (ex: **[CLS]** start of a sample, **[SEP]** sentences separator, etc.)
- Solution:
  - Align NER tag labels to sentence tokens**

Tokens:	EU	rejects	German	call	to	boycott	British	lamb	.
NER tag no.:	3	0	7	0	0	0	7	0	0



str tokens:	[CLS]	EU	rejects	German	call	to	boycott	British	la	##mb	.	[SEP]
int tokens:	101	7270	22961	1528	1840	1106	21423	1418	2495	12913	119	102
aligned tags:	-100	3	0	7	0	0	0	7	0	0	0	-100

example of **NER tag alignment**: 1<sup>st</sup> sentence of CoNLL2003 NER tags aligned

## Step 1: Create Tokenized Dataset - steps

- Taken together, the tokenized dataset was created in the following three steps

```
def align_labels_with_tokens(labels, word_ids):
    new_labels = []
    current_word = None
    for word_id in word_ids:
        if word_id != current_word:
            current_word = word_id
            label = -100 if word_id is None else labels[word_id]
            new_labels.append(label)
        elif word_id is None:
            new_labels.append(-100)
        else:
            label = labels[word_id]

    # print(f'before: {label}')
    if label % 2 == 1:
        label += 1
    # print(f'after: {label}')
    new_labels.append(label)
    return new_labels
```



```
def tokenize_and_align_labels(examples):
    tokenized_inputs = tokenizer(
        examples['tokens'], truncation = True, is_split_into_words = True
    )
    all_labels = examples['ner_tags']
    new_labels = []
    for i, labels in enumerate(all_labels):
        word_ids = tokenized_inputs.word_ids(i)
        new_labels.append(align_labels_with_tokens(labels, word_ids))
    tokenized_inputs['labels'] = new_labels
    return tokenized_inputs
```



```
tokenized_dataset = dataset.map(
    tokenize_and_align_labels,
    batched = True,
    remove_columns = dataset['train'].column_names
)
```

## Step 2: Create Input Objects for Trainer

- **Tokenizer**
- **Evaluation metrics function**
- **Data collator**

```
from transformers import AutoTokenizer  
tokenizer = AutoTokenizer.from_pretrained("bert-base-cased")
```

```
def compute_metrics(eval_preds):  
    logits, labels = eval_preds  
    predictions = np.argmax(logits, axis=2)  
  
    true_labels = [[label_names[l] for l in label if l != -100] for label in labels]  
    true_predictions = [  
        [label_names[p] for (p, l) in zip(prediction, label) if l != -100]  
        for prediction, label in zip(predictions, labels)  
    ]  
    all_metrics = metric.compute(predictions=true_predictions, references=true_labels)  
    return {  
        "precision": all_metrics["overall_precision"],  
        "recall": all_metrics["overall_recall"],  
        "f1": all_metrics["overall_f1"],  
        "accuracy": all_metrics["overall_accuracy"],  
    }
```

```
from transformers import DataCollatorForTokenClassification  
data_collator = DataCollatorForTokenClassification(tokenizer = tokenizer)
```

## Step 3: Define BERT Model and Training Arguments

---

- Pre-trained **BERT** model

```
from transformers import AutoModelForTokenClassification
model = AutoModelForTokenClassification.from_pretrained(
    'bert-base-cased',
    num_labels = 9,
)
```

- Training arguments

```
# specify training arguments to pass to the trainer
training_args = TrainingArguments(
    output_dir="bert-conll2003-test-output",
    evaluation_strategy="epoch",
    per_device_train_batch_size=16,
    per_device_eval_batch_size=16,
    num_train_epochs=3,
    learning_rate=2e-5,
    weight_decay=0.01
)
```

## Step 4: Train and evaluate the BERT Model on CoNLL2003

---

- Put the objects created in the previous three steps together and create the following trainer. **Train** and **evaluate** the BERT model

```
from transformers import Trainer
trainer = Trainer(
    model = model,
    args = args,
    train_dataset = tokenized_dataset['train'],
    eval_dataset = tokenized_dataset['validation'],
    data_collator = data_collator,
    compute_metrics = compute_metrics,
    tokenizer = tokenizer,
)

trainer.train()
```



## Evaluate the Training Results of the BERT Model

---

- Final configuration
  - *compute\_metrics* function: *axis* = 2
  - *args*: *num\_train\_epochs* = 3, *per\_device\_train\_batch\_size* = 16, *per\_device\_eval\_batch\_size* = 16, *learning\_rate* = 2e-5, *weight\_decay*=0.01

Epoch	Eval Loss	Eval Precision	Eval Recall	Eval F1	Eval Accuracy
1	0.0416	0.9287	0.9384	0.9335	0.9895
2	0.0392	0.9402	0.9473	0.9438	0.9907
3	0.0365	0.9444	0.9519	0.9481	0.9912

## Conclusions

---

- *WikiNER* brutal force model is **ineffective** to perform a NER task
  - **Low performances** (max recall: 29% - 10,000 NEs and aliases)
  - **Time-consuming** (it takes hours)
  - **No NEs types**
- Two issues in *Wikipedia* entities:
  - The presence of **more NEs** naturally **decreases precision** (although **increases recall**)
  - The *Wikipedia* entities are **not universal** enough to capture benchmark entities
- Given proper configurations in *TrainingArguments()*, **BERT** model **performs notably better** than the brute-force model, identifying **NE type** too
  - The greater the batch size, the greater the performances

Questions?

**Thank You!**